

Technical Assignment — Gates Foundation AI Fellowship – India 2026

Overview

This assignment is part of the technical screening process for the fellowship. Candidates are asked to install, assess, and build on a real AI evaluation tool. We are interested both in the technical output and in the decisions made to produce it.

Submission of a code repository is a prerequisite for further consideration. Candidates who do not submit a repository will not proceed.

The Tool

Install the CeRAI AI Evaluation Tool: <https://github.com/cerai-iitm/AIEvaluationTool>

This is an open-source framework developed at CeRAI, IIT Madras, for the automated evaluation of conversational AI systems. It evaluates conversational endpoints by sending structured text inputs and assessing responses against quality metrics covering accuracy, safety, and user experience. It supports API endpoints, WhatsApp bots, and web-based interfaces.

Candidates are expected to read the documentation carefully before proceeding. Understanding the scope and limitations of the tool is itself part of the task.

Choose One Path

Option A — Evaluate & Report

Select a conversational endpoint of your choice, design a test suite, and run it using the tool. Your repository should contain all code written or materially modified, with a README sufficient for another developer to reproduce your setup without needing to contact you.

Your live endpoint should present your findings as a self-contained document. A strong submission will include: a) what conversational system was evaluated and why it was chosen, b) what test suite was designed and the reasoning behind it, c) what the results show with explicit interpretation rather than raw data alone, d) what conclusions can be drawn, and e) where the findings fall short or should not be generalized. Where possible, embed a structured data block summarizing the key findings in machine-readable form.

Option B — Critique & Rebuild

If you determine that the tool is fundamentally limited or unsuitable for evaluating conversational endpoints, make the case systematically. File detailed issues on the repository documenting what you found — each issue should include a description of the problem, steps to reproduce it, its impact on evaluation quality, and a suggested fix. Then implement a minimal viable alternative capable of sending text inputs to a conversational endpoint and assessing the responses. This may be done using other existing open source frameworks.

Your repository should contain the alternative implementation (or steps to reproduce if using an existing open source framework) and a README sufficient for another developer to get it running. Your live endpoint should present both the critique and the alternative as a self-contained document. A strong submission will include: what specifically is broken or insufficient in the existing tool with reference to the issues filed, what design decisions were made in implementing the alternative and why, a demonstration that the alternative works against a conversational endpoint of your choice, and an honest account of what the alternative still does not do well.

In either case, your choice of path and the reasoning behind it may be discussed at a later stage.

How to Submit

Your submission should include the following:

1. Your repository URL (GitHub or equivalent, set to public)
2. Your live endpoint URL
3. A single paragraph describing the path you chose and why
4. A description of AI use in completing the assignment - what approach did you take, where did you have to course correct, how did you do so?

Submissions that do not include a repository URL will not be reviewed. If your endpoint is not yet live at time of submission, include it in the repository as a documented deliverable with instructions to run it locally.

Partial submissions are accepted. A repository with clear documentation of what was attempted and where it broke down is preferable to no submission.

Please share back to Drew Durlofsky within 2 days of receipt.